

# Vector Load Balancing in Charm++

---

Ronak Buch

October 19, 2022

Parallel Programming Laboratory  
University of Illinois at Urbana-Champaign


# Background and Motivation

---


# Dynamic Load Balancing

- Adaptively rearrange work amongst PEs to maximize performance as computation evolves
- Most loaded PE usually determines iteration time
- Enabled by migratable objects, runtime instrumentation, overdecomposition
- Necessary for scaling all but very regular, static applications

# Before LB



# After LB



# Scalar Load

- To measure load traditionally:
  - Start timer when object begins execution
  - End timer when control returns to RTS
  - Elapsed time added to object's load
  - PE's load is sum of resident objects' loads

# Scalar Load

- To measure load traditionally:
  - Start timer when object begins execution
  - End timer when control returns to RTS
  - Elapsed time added to object's load
  - PE's load is sum of resident objects' loads
- Only captures time object spends active
  - CPU time alone does not determine performance
  - Deficient for many classes of applications

# Scalar Load Deficiencies

- Phase-based applications
  - Iteration divided into several orthogonal phases
  - Time spent in different phases not fungible
  - Scalar loses temporal execution constraints



# Scalar Load Deficiencies

- Phase-based applications
  - Iteration divided into several orthogonal phases
  - Time spent in different phases not fungible
  - Scalar loses temporal execution constraints
- Resource constrained applications
  - Memory footprint, FD limit, cache working set size, ...

# Scalar Load Deficiencies

- Phase-based applications
  - Iteration divided into several orthogonal phases
  - Time spent in different phases not fungible
  - Scalar loses temporal execution constraints
- Resource constrained applications
  - Memory footprint, FD limit, cache working set size, ...
- Asynchronous computation
  - Object invokes and waits for GPU kernel, async I/O, ...

# Scalar Load Deficiencies

- Phase-based applications
  - Iteration divided into several orthogonal phases
  - Time spent in different phases not fungible
  - Scalar loses temporal execution constraints
- Resource constrained applications
  - Memory footprint, FD limit, cache working set size, ...
- Asynchronous computation
  - Object invokes and waits for GPU kernel, async I/O, ...
- Cannot be captured in single scalar value!

# What is Load?

- Load is a quantification of the characteristics of entities in an optimization problem

# What is Load?

- Load is a quantification of the characteristics of entities in an optimization problem
- Load often consists of more than one metric:
  - In shipping: consider cargo volume and weight
  - In nutrition: consider fat, protein, carbohydrates, ...

# What is Load?

- Load is a quantification of the characteristics of entities in an optimization problem
- Load often consists of more than one metric:
  - In shipping: consider cargo volume and weight
  - In nutrition: consider fat, protein, carbohydrates, ...
- The same applies to computer programs!

# Vector Load Balancing

---

# Vector Load

- Instead of using a single scalar  $l$  for load, use a vector  $\vec{l} = \langle l_1, l_2, \dots, l_d \rangle$  of dimension  $d$




# Vector Load

- Instead of a using single scalar  $l$  for load, use a vector  $\vec{l} = \langle l_1, l_2, \dots, l_d \rangle$  of dimension  $d$
- Remedies earlier deficiencies:
  - Phase-based applications
    - Time spent in each phase:  $\langle t_A, t_B, \dots, t_n \rangle$
  - Resource constrained applications
    - Resource usage alongside CPU time:  $\langle cpu, mem \rangle$
  - Asynchronous computation
    - GPU time alongside CPU time:  $\langle cpu, gpu \rangle$
- Rich, flexible, and extensible packaging of load


# Measuring Vector Loads

- Features and APIs to measure vector loads:
  - Applications can call function to indicate phase boundary, RTS automatically measures per-phase load
  - Runtime flags to automatically add communication load (msgs, bytes sent)
  - Memory footprint via PUP
  - GPU load via `acce1` or CUDA timers
  - Users may manually specify load vector


# Phase-Based Application



# Phase Unaware LB



# Phase Aware LB



# Vector Load Balancing Strategies

---

# Vector Balancing

- Existing LB strategies cannot use a load vector
  - Still compatible, vector converted to scalar via sum, max, etc.
- Multiple dimensions makes vector load balancing more complex
  - Objects can no longer be totally ordered
  - Want to minimize over all dimensions simultaneously
    - Single variable optimization is now multivariate
- New LB strategies are needed

# Vector Strategies - Greedy

- Extension of scalar greedy strategy to vector loads
  - Scalar version goes through objects from heaviest to lightest and assigns to the current least loaded processor



# Vector Strategies - Greedy

- Extension of scalar greedy strategy to vector loads
  - Scalar version goes through objects from heaviest to lightest and assigns to the current least loaded processor
- Vector version:
  - Create  $d$  PE minheaps, each keyed on a different dimension of the vector, add all PEs to each heap
  - Go through objects in descending  $\max(\vec{l})$ , assign to minimum PE in dimension of  $\max(\vec{l})$  and update heaps
- Simple, but focuses on single dimension at a time

# Vector Strategies - METIS

- METIS supports giving vertices vector weights
- Reframe LB problem as graph, objects map to vertices and output partitions map to PEs
  - No edges, but could use with comm graph
- Implemented via bipartitioning objects based on  $\max(\vec{l})$  like Greedy, but adds extra refinement phase
  - Graph coarsening/refinement for sake of performance
- Generally works pretty well, but gives poor results for some configurations

# Vector Strategies - Norm

- Go through objects in descending  $\max(\vec{l})$  and place on the PE such that the post-placement PE load vector norm is globally minimized
  - Works well, but computationally expensive
  - Norm inequalities are not preserved under vector addition:
    - $\|(2, 0)\|_2 < \|(0, 3)\|_2$
    - $\|(2, 0) + (3, 0)\|_2 > \|(0, 3) + (3, 0)\|_2$
  - Makes it non-trivial to reduce search space
- Choice of underlying norm (2-norm,  $\infty$ -norm, etc.)

# Vector Strategies - Norm

- Implemented several different versions of norm-based assignment
  - Exhaustive search
  - $k$ -d tree-based search
  - $rk$ -d tree-based search
  - Pareto frontier-based search
- All give same result, but different performance

# Vector Strategies - Norm

- Implemented several different versions of norm-based assignment
  - Exhaustive search
  - $k$ -d tree-based search
  - $rk$ -d tree-based search
  - Pareto frontier-based search
- All give same result, but different performance
- The  $rk$ -d variant generally performs best

# NormLB - $k$ -d

- Represent PEs as points in a  $k$ -d tree
  - Arbitrary dimensional space partitioning tree
  - Can prune search space as candidates are found
- $k$ -d works well for searching in static point set, but here, tree updated after every assignment
  - Costly update operations
  - Structured pattern of updates results in unbalanced tree
- Can be worse than the naïve exhaustive version!

# Random Relaxed $k$ -d

- **Random Relaxed**  $k$ -d trees help solve these problems; two key differences from standard  $k$ -d:
  - Random** Discriminant is uniformly randomly chosen and each insertion has some probability of becoming the root, or root of subtree, ...
  - Relaxed** Instead of cycling through discriminants,  $1, 2, \dots, k, 1, \dots$ , each node stores arbitrary discriminant  $j \in \{1, 2, \dots, k\}$
- Stochasticity keeps tree balanced, makes updates fast

# Random Relaxed $k$ -d




Figure 1:  $k$ -d





Figure 2:  $rk$ -d



# $k$ -d - Pruning



# NormLB - Performance

Method	Strategy Time (s)	
	<i>1e4 PEs, 1e5 objs</i>	<i>1e4 PEs, 1e6 objs</i>
Exhaustive	2.18	21.54
Standard $k$ -d	0.93	27.55
Relaxed $k$ -d	0.57	7.96

**Table 1:** Performance of Norm-Based Strategies

Synthetic data: 2 phase (exp  $\lambda = 0.15$ , normal  $\mu = 10, \sigma^2 = 3$ )

# Objectives and Results

---

# Phase Objective Function

$M$  mapping,  $O$  set of objects,  $P$  set of PEs

$$\arg \min_M \left( \sum_{0 \leq i < d} \left( \max_{p \in P} \left( \sum_{\forall o \in O: M(o)=p} (\vec{l}_o)_i \right) \right) \right) \quad (1)$$

# Phase Objective Function

$M$  mapping,  $O$  set of objects,  $P$  set of PEs

$$\arg \min_M \left( \sum_{0 \leq i < d} \left( \max_{p \in P} \left( \sum_{\forall o \in O: M(o)=p} (\vec{l}_o)_i \right) \right) \right) \quad (1)$$

Load on PE  $p$  in dimension  $i$

# Phase Objective Function

$M$  mapping,  $O$  set of objects,  $P$  set of PEs

$$\arg \min_M \sum_{0 \leq i < d} \left( \max_{p \in P} \left( \sum_{\forall o \in O: M(o)=p} (\vec{l}_o)_i \right) \right) \quad (1)$$

Load on PE  $p$  in dimension  $i$

Maximum load in dimension  $i$  on a single PE

# Phase Objective Function

$M$  mapping,  $O$  set of objects,  $P$  set of PEs

Sum across dimensions

Load on PE  $p$  in dimension  $i$

$$\arg \min_M \left( \sum_{0 \leq i < d} \left( \max_{p \in P} \left( \sum_{\forall o \in O: M(o)=p} (\vec{l}_o)_i \right) \right) \right) \quad (1)$$

Maximum load in dimension  $i$  on a single PE

# Overlapped Objective Function

$M$  mapping,  $O$  set of objects,  $P$  set of PEs

$$\arg \min_M \max_{0 \leq i < d} \left( \max_{p \in P} \left( \sum_{\forall o \in O: M(o)=p} (\vec{l}_o)_i \right) \right) \quad (2)$$

Load on PE  $p$  in dimension  $i$

Maximum load in dimension  $i$  on a single PE



# Overlapped Objective Function

$M$  mapping,  $O$  set of objects,  $P$  set of PEs

Max across dimensions

Load on PE  $p$  in dimension  $i$

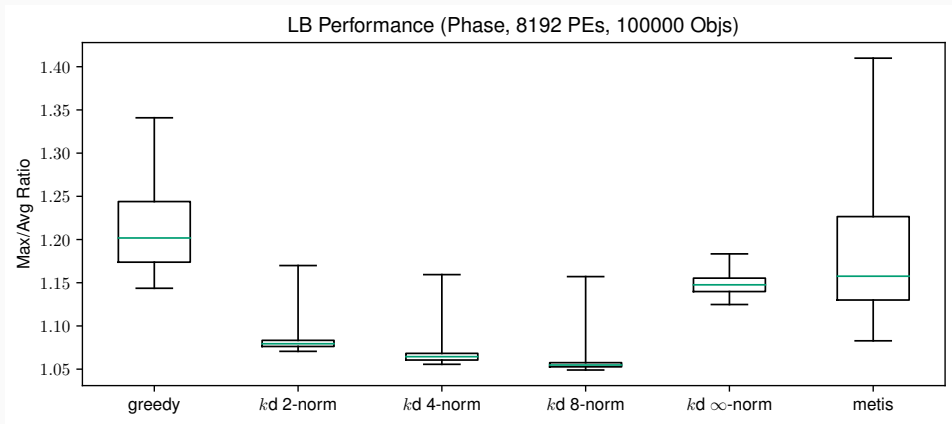
$$\arg \min_M \left( \max_{0 \leq i < d} \left( \max_{p \in P} \left( \sum_{\forall o \in O: M(o)=p} (\vec{l}_o)_i \right) \right) \right) \quad (2)$$

Maximum load in dimension  $i$  on a single PE

The diagram illustrates the components of the objective function equation (2). The equation is: 
$$\arg \min_M \left( \max_{0 \leq i < d} \left( \max_{p \in P} \left( \sum_{\forall o \in O: M(o)=p} (\vec{l}_o)_i \right) \right) \right) \quad (2)$$
 Annotations include: 


- A green arrow labeled "Max across dimensions" pointing to the  $\max_{0 \leq i < d}$  term.
- A blue arrow labeled "Load on PE  $p$  in dimension  $i$ " pointing to the sum  $\sum_{\forall o \in O: M(o)=p} (\vec{l}_o)_i$ .
- A red arrow labeled "Maximum load in dimension  $i$  on a single PE" pointing to the  $\max_{p \in P}$  term.

# Vector LB Simulations - Phase



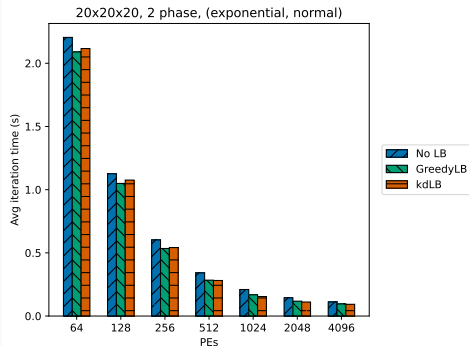
100 trials, Synthetic data: 2 phase (exp  $\lambda = 0.15$ , normal  $\mu = 10, \sigma^2 = 3$ )

# Vector LB Simulations - Overlapped

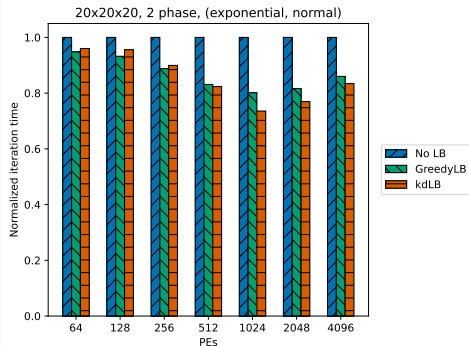


100 trials, Synthetic data: 2 phase (exp  $\lambda = 0.15$ , normal  $\mu = 10, \sigma^2 = 3$ )

# Vector LB Runs



Runtime



Normalized

Results from KNL partition of Stampede2, 2 phase (exp  $\lambda = 0.15$ , normal  $\mu = 10, \sigma^2 = 3$ )

- Task-based programming model from Sandia
- Collaborated with team to create two adapters:
  1. To allow Charm++ LBs to be used in VT
  2. To ingest VT logs into Charm++ LB simulator
- Phase-based application called EMPIRE
  - 14 dimensions
- *kdLB* gives approximately 12% performance improvement over previous best LB strategy




Figure 3: With previous best LB

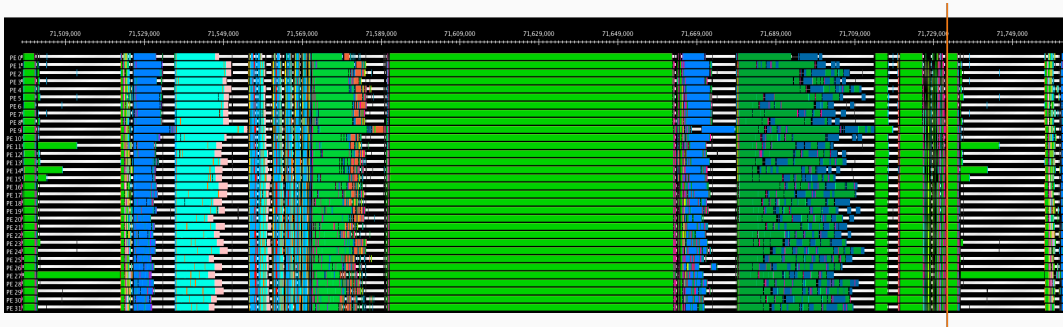


Figure 4: With  $k$ dLB

# Additional Topics

- *kdConstrainLB* allows subset of dimensions to be constrained while rest are minimized
  - Prevented `malloc` failures in memory constrained run
  - Other LB strategies resulted in crashes
- Additional class of approximate norm-based strategies that tradeoff quality for performance



# Conclusions

- Load characteristics of complex, modern applications cannot be captured in a single scalar
- New LB strategies can tractably utilize the additional detail provided by a load vector
- Vector LB has shown improvement over scalar LB across synthetic and production applications

**Questions?**